

Junjie Yu

Income Level Prediction

Overview

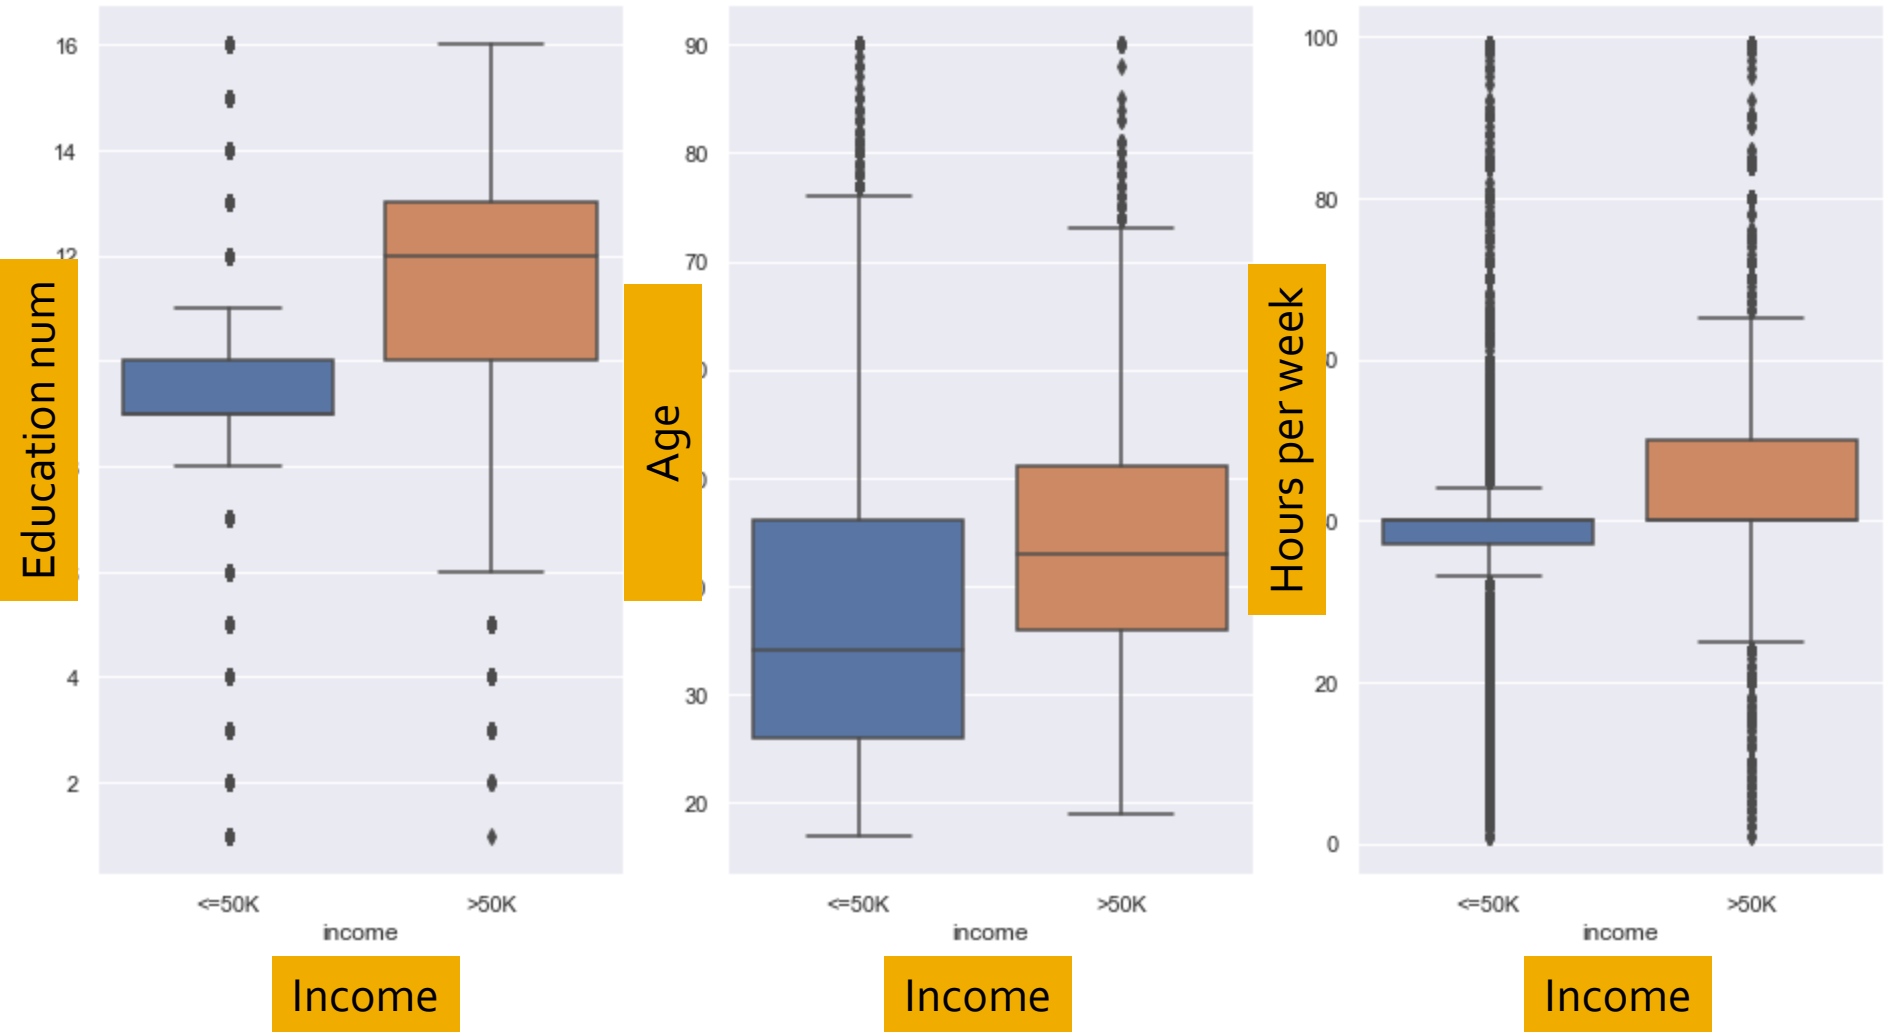
- Objective:
 - To predict whether a person makes over 50K a year.
- Data:
 - 1994 US census database
 - ~ 32,000 observations with 14 variables
 - Source: <http://archive.ics.uci.edu/ml/datasets/Adult>

Data


- 13 features
- 1 target variable: income ($\leq 50k$, $> 50k$)

Age	Workclass	Education level	Education-num	Marital-status	Occupation	Relationship	Race	Sex	Capital-gain	Capital-loss	Hours/week	Native-country	Income
int	object	object	float	object	object	object	object	object	float	float	float	object	object
39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	$\leq 50K$
50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	$\leq 50K$
38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	$\leq 50K$
53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	$\leq 50K$
28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	$\leq 50K$


EDA: Education, age, and hours per work



EDA: Education num and education level




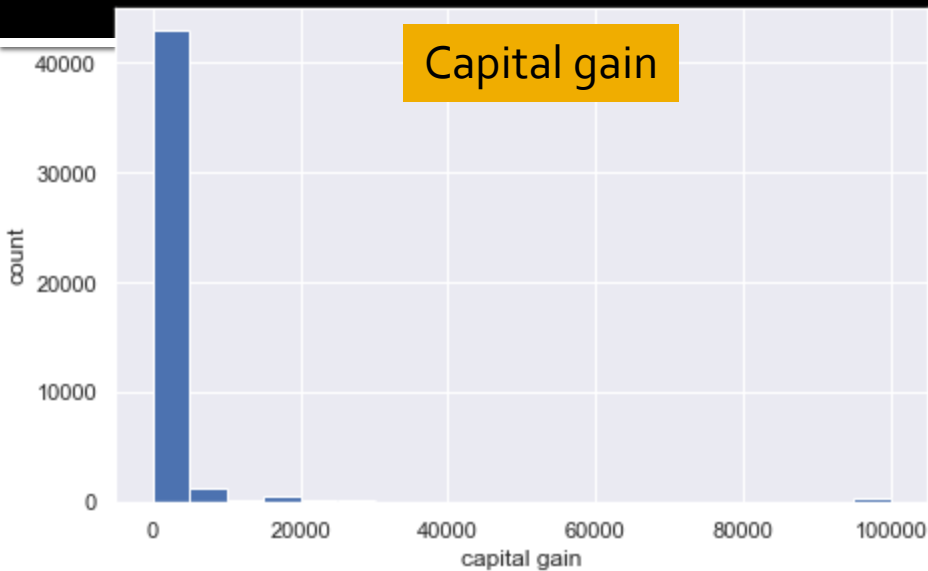


Education num




Education level


EDA: Capital gain and capital loss



EDA: Occupation and native country




Occupation



Native country

EDA: Marital status and relationship






Marital status

Relationship

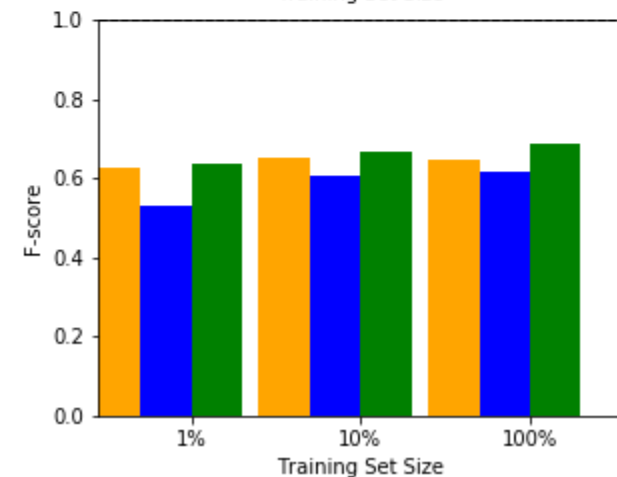


Gradient Boosting is Selected

GaussianNB DecisionTreeClassifier GradientBoostingClassifier

Training data



Validation data



Feature Engineering

Initial

- Initial model with raw features

Education

- Discard redundant feature education level

Capital

- Combine capital gain and capital loss

Native


- Recode native country as U.S. and other

Optimize

- Optimize hyper-parameters


Model Evolution

Accuracy	0.8630
Precision	0.7821
Recall	0.6073




Model Evolution

Accuracy	0.8636
Precision	0.7831
Recall	0.6091




Model Evolution

Accuracy	0.8636
Precision	0.7831
Recall	0.6091




Model Evolution

Accuracy	0.8640
Precision	0.7849
Recall	0.6091



Model Evolution

Accuracy	0.8714
Precision	0.7853
Recall	0.6503



Summary

- Gradient Boosting is selected over Gaussian Naïve Bayes and Decision Tree based on running time, accuracy, and F score.
- Model is improved/simplified progressively by feature engineering:
 - Discarding redundant feature
 - Combining capital gain & capital loss
 - Recoding native country